

# Dataset Description

CRIM: Per capita crime rate by town. ZN: Proportion of residential land zoned for lots over 25,000 sq. ft. INDUS: Proportion of non-retail business acres per town. CHAS: Charles River dummy variable (1 if tract bounds river; 0 otherwise). NOX: Nitric oxides concentration (parts per 10 million). RM: Average number of rooms per dwelling. AGE: Proportion of owner-occupied units built prior to 1940. DIS: Weighted distances to five Boston employment centers. RAD: Index of accessibility to radial highways. TAX: Full-value property tax rate per 10,000. *PTRATIO* : *Pupil – teacherratio by town*.  $B : 1000(B_k - 0.63)^2$  where  $B_k$  is *MEDV* : *Median value of owner – occupied homes in 1000s* (this is the target variable)

```
In [ ]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score
```

```
In [ ]: df=pd.read_csv("HousingData.csv")
```

```
In [ ]: df.head()
```

Out [ ]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B I
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1	296	15.3	396.90
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2	242	17.8	396.90
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2	242	17.8	392.83
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3	222	18.7	394.63
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3	222	18.7	396.90

```
In [ ]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 506 entries, 0 to 505
Data columns (total 14 columns):
 #   Column      Non-Null Count  Dtype  
---  -
 0   CRIM        486 non-null    float64
 1   ZN          486 non-null    float64
 2   INDUS       486 non-null    float64
 3   CHAS        486 non-null    float64
 4   NOX         506 non-null    float64
 5   RM          506 non-null    float64
 6   AGE         486 non-null    float64
 7   DIS         506 non-null    float64
 8   RAD         506 non-null    int64   
 9   TAX         506 non-null    int64   
10  PTRATIO     506 non-null    float64
11  B           506 non-null    float64
12  LSTAT       486 non-null    float64
13  MEDV        506 non-null    float64
dtypes: float64(12), int64(2)
memory usage: 55.5 KB
```

```
In [ ]: df.isnull().sum()
```

```
Out[ ]: CRIM      20
        ZN      20
        INDUS  20
        CHAS   20
        NOX     0
        RM     0
        AGE    20
        DIS     0
        RAD     0
        TAX     0
        PTRATIO 0
        B       0
        LSTAT   20
        MEDV    0
dtype: int64
```

```
In [ ]: df=df.fillna(df.mean())
```

```
In [ ]: X = df.iloc[:, 0:-1]
        y = df.iloc[:, -1]
```

```
In [ ]:
```

```
In [ ]: print(X)
```

Reg										
	CRIM	ZN	INDUS	CHAS	NOX	RM	\			
0	0.00632	18.0	2.31	0.0	0.538	6.575				
1	0.02731	0.0	7.07	0.0	0.469	6.421				
2	0.02729	0.0	7.07	0.0	0.469	7.185				
3	0.03237	0.0	2.18	0.0	0.458	6.998				
4	0.06905	0.0	2.18	0.0	0.458	7.147				
..	...	...	...	...	...	...				
501	0.06263	0.0	11.93	0.0	0.573	6.593				
502	0.04527	0.0	11.93	0.0	0.573	6.120				
503	0.06076	0.0	11.93	0.0	0.573	6.976				
504	0.10959	0.0	11.93	0.0	0.573	6.794				
505	0.04741	0.0	11.93	0.0	0.573	6.030				
							AGE	DIS	RAD	TAX \
0							65.2	4.0900	1	296
1							78.9	4.9671	2	242
2							61.1	4.9671	2	242
3							45.8	6.0622	3	222
4							54.2	6.0622	3	222
..							...	...	...	...
501							69.1	2.4786	1	273
502							76.7	2.2875	1	273
503							91.0	2.1675	1	273
504							89.3	2.3889	1	273
505	<bound method DataFrame.mean of						CRIM	...	2.5050	1 273
	PTRATIO	B							LSTAT	
0	15.3	396.90							4.98	
1	17.8	396.90							9.14	
2	17.8	392.83							4.03	
3	18.7	394.63							2.94	
4	18.7	396.90	<bound method DataFrame.mean of						CRIM ...	
..	...	...	<bound method DataFrame.mean of						CRIM ...	
501	21.0	391.99	<bound method DataFrame.mean of						CRIM ...	
502	21.0	396.90							9.08	
503	21.0	396.90							5.64	
504	21.0	393.45							6.48	
505	21.0	396.90							7.88	

[506 rows x 13 columns]

```
In [ ]: print(y)

0      24.0
1      21.6
2      34.7
3      33.4
4      36.2
...
501     22.4
502     20.6
503     23.9
504     22.0
505     11.9
Name: MEDV, Length: 506, dtype: float64

In [ ]: X_train,X_test,y_train,y_test=train_test_split(X,y,random_state=42,test_size=0.2)

In [ ]: X_test=X_test.apply(pd.to_numeric,errors='coerce')

In [ ]: X_train = X_train.apply(pd.to_numeric, errors='coerce')
```

```
In [ ]: X_train = X_train.astype(float)
        y_train = y_train.astype(float)
```

```
In [ ]: Lr=LinearRegression()
```

```
In [ ]: len(X_train)
```

```
Out[ ]: 354
```

```
In [ ]: print(X_train)
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX
5	0.02985	0.0	2.18	0.0	0.458	6.430	58.7	6.0622	3.0	222.0
116	0.13158	0.0	10.01	0.0	0.547	6.176	72.5	2.7301	6.0	432.0
45	0.17142	0.0	6.91	0.0	0.448	5.682	33.8	5.1004	3.0	233.0
16	1.05393	0.0	8.14	0.0	0.538	5.935	29.3	4.4986	4.0	307.0
468	15.57570	0.0	18.10	0.0	0.580	5.926	71.0	2.9084	24.0	666.0
...	...	...	...	...	...	...	...	...	...	...
106	0.17120	0.0	8.56	0.0	0.520	5.836	91.9	2.2110	5.0	384.0
270	0.29916	20.0	6.96	0.0	0.464	5.856	42.1	4.4290	3.0	223.0
348	0.01501	80.0	2.01	0.0	0.435	6.635	29.7	8.3440	4.0	280.0
435	11.16040	0.0	18.10	0.0	0.740	6.629	94.6	2.1247	24.0	666.0
102	0.22876	0.0	8.56	0.0	0.520	6.405	85.4	2.7147	5.0	384.0
...	...	...	...	...	...	...	...	...	...	...
5	18.7	394.12	5.21							
116	17.8	393.30	NaN							
45	17.9	396.90	10.21							
16	21.0	386.85	6.58							
468	20.2	368.74	18.13							
...	...	...	...							
106	20.9	395.67	18.66							
270	18.6	388.65	13.00							
348	17.0	390.94	5.99							
435	20.2	109.85	23.27							
102	20.9	70.80	10.63							

[354 rows x 13 columns]

```
In [ ]: X_train = X_train.dropna()
        y_train = y_train[X_train.index]
        X_test=X_test.dropna()
        y_test=y_test[X_test.index]
```

```
In [ ]: print(y_train)
```

```
5      28.7
45     19.3
16     23.1
468    19.1
360    25.0
...
106    19.5
270    21.1
348    24.5
435    13.4
102    18.6
Name: MEDV, Length: 267, dtype: float64
```

```
In [ ]: len(y_train)
```

Out[ ]: 267

```
In [ ]: Lr.fit(X_train,y_train)
```

```
Out[ ]: LinearRegression
LinearRegression()
```

```
In [ ]: pred=Lr.predict(X_test)
```

```
In [ ]: print(len(pred))


127
```

```
In [ ]: #Checking Accuracy Score
print('Test Score',Lr.score(X_test,y_test))
print('Train Score',Lr.score(X_train,y_train))
```

Test Score 0.7197008120278565  
Train Score 0.7786240464453613

```
In [ ]: import seaborn as sns
sns.displot(pred-y_test, kind = 'kde')
```

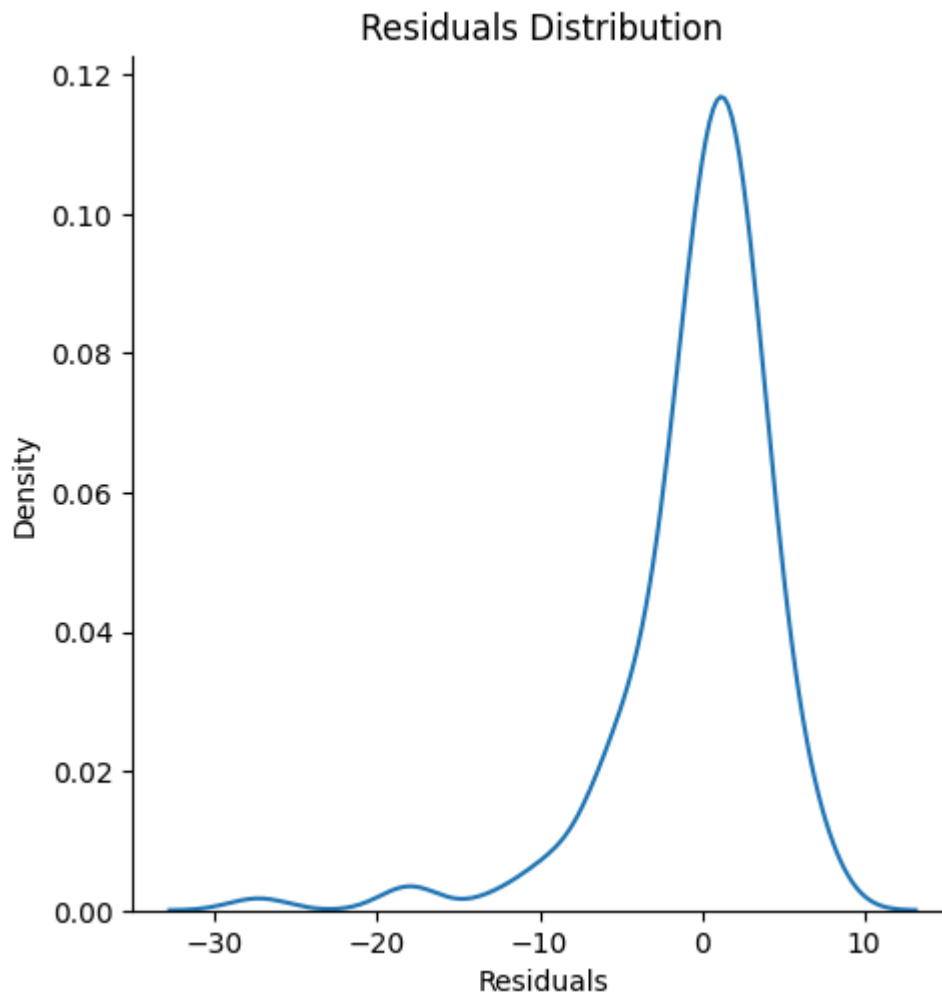
Out[ ]: <seaborn.axisgrid.FacetGrid at 0x30e053250>



```
In [ ]: from sklearn.metrics import r2_score
score = r2_score(pred, y_test)
score
```

Out[ ]: 0.6603974893117284

```
In [ ]: # Plot residuals
sns.displot(pred - y_test, kind='kde')
plt.title('Residuals Distribution')
plt.xlabel('Residuals')
plt.ylabel('Density')
plt.show()
```



```
In [ ]: # Plot predicted vs actual values
plt.figure(figsize=(10, 6))
plt.scatter(y_test, pred, alpha=0.6)
plt.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()], 'r--',
plt.xlabel('Actual MEDV')
plt.ylabel('Predicted MEDV')
plt.title('Predicted vs. Actual MEDV')
plt.show()
```

